

A1 – Opinion Ordering

Junior judges vote first. The rule exists to prevent anchoring: a reasoning agent that hears an authoritative opinion before forming its own cannot later claim to have reasoned independently, even if it arrives at the same conclusion.

THE ALGORITHM

```
[
INPUT: panel[], query q

for judge in panel.order_by(seniority ASC):  // junior first
    opinion = judge.generate(q, prior_opinions=[])  // no exposure to prior
    opinions.append(opinion)

return opinions]
```

Constraint: Each judge generates their opinion with `prior_opinions=[]`. No judge sees any prior vote before submitting their own.

Rationale for the ordering: Reversing the loop – senior first – does not merely reorder the output. It changes what the output represents. A junior judge who hears a senior opinion before voting is no longer producing independent reasoning; they are producing a response to an authoritative prior. The outputs look the same. The process is categorically different.

THE HALACHIC SOURCE

Mishnah Sanhedrin 4:2: Deliberation begins from the side seats. The least senior judges are called upon first.

Bartenura's rationale (15th century): The rule is not about humility or procedure – it is about epistemic contamination. A junior judge who hears the senior opinion before forming his own has been contaminated regardless of whether he would have reached the same conclusion independently. The contamination is invisible in the output and only exists in the process. This is why the rule is structural: there is no way to "account for"

anchoring after the fact. The only solution is to prevent exposure before commitment.

Why this matters: Bartenura's framing – contamination is invisible in output, visible only in process – is the insight that makes A1 relevant to AI training. Evaluating model outputs for sycophancy is insufficient. A model that produced the correct answer via sycophantic capitulation is a training failure even if the answer is right.

THE ALIGNMENT APPLICATION

The mapping is direct. A1 is the procedural implementation of P2's Phase 1 requirement:

- **Seniority ordering** → **authority signals**: Senior judges = high-confidence prior responses, authoritative sources, or system context that the model should not anchor to before forming an initial position.
- **prior_opinions=[]** → **clean Phase 1**: The model generates its initial position without exposure to what others have concluded.
- **Junior-first** → **commitment before context**: The model commits to a position before receiving social signals. It can then update – but the before/after is always visible.

The loop structure in A1 also explains why **multi-agent deliberation is architecturally superior to single-agent reasoning** for high-stakes decisions: independent opinions generated without cross-contamination produce genuine signal variance. A panel where opinions are generated sequentially without prior exposure will produce the right distribution of disagreements. A panel where agents see each other's outputs first will converge artificially.

DERIVED PRINCIPLE

A1 is the procedural foundation of **P2 (Anti-Sycophancy: Independent Reasoning First)**.

P2 is the principle; A1 is one implementation. The Phase 1/Phase 2 structure of P2 is A1 applied to a single agent across time: Phase 1 is the agent's `prior_opinions=[]` moment, Phase 2 is after commitment.

A1 also informs the **HMTF Turn 0 anchor mechanism** – the anchor stored outside conversation history is the Phase 1 output that subsequent turns cannot overwrite.

SEE ALSO

- [P2 – Anti-Sycophancy: Independent Reasoning First](#) – the principle A1 grounds
- [A2 – Unanimous Conviction Rule](#) – what to do after A1 produces unanimous output
- [A5 – Judge Eligibility Check](#) – must run before A1
- [HMTP – Halachic Multi-Turn Protocol](#) – A1 applied across conversation turns
- [knowledge/claude/algorithms.md](#) – lean pseudocode reference