A5 — Judge Eligibility Check

Disqualification is structural, not self-assessed. A judge who has a financial interest, a relationship to a party, or a prior public position on the exact question does not assess whether they can be objective — they are disqualified. The disqualification is binary and runs before deliberation begins.

THE ALGORITHM

```
INPUT: judge j, query q

DISQUALIFY(j) IF any of:
    j.financial_interest_in(q) == True
    j.related_to_party(q) == True
    j.prior_ruling_in_conflicting_capacity(q) == True
    j.public_prior_position_on(q) == True

IF disqualified:
    j.recuse()
    j.disclose_reason()
    RETURN: cannot participate in this case

// Passed — eligible to proceed
RETURN: eligible

TIMING: MUST run before A1. Non-negotiable.
```

The disclosure is mandatory. A disqualified judge does not simply step aside silently. They disclose the structural reason for their disqualification. This makes the disqualification auditable and prevents the reason from being unknown to the parties.

The check is structural, not subjective. The algorithm does not call `j.believes_they_can_be_objective()`. That function is not in scope. The question asked is: "Does j have a structural relationship that a reasonable external observer would identify as a conflict?" If yes, disqualified.

THE HALACHIC SOURCE

Rambam, Hilchot Sanhedrin 2:1–3: A judge related to either party in any degree covered by the enumerated relationships is disqualified. A judge with a financial stake in the outcome is disqualified. A judge who has previously ruled on the same matter in a conflicting capacity is disqualified.

The key legal principle: The standard for disqualification is not the judge's subjective confidence in their own objectivity — it is the existence of a structural reason that an external observer would recognize as a bias risk. Rambam explicitly rejects the self-assessment model: a judge who believes they can be fair is not thereby qualified. The structural relationship is the deciding criterion.

Why this matters: Self-assessed bias monitoring is the default in both judicial systems and AI alignment frameworks. The problem is that the instrument doing the assessment (the judge, the model) is exactly the instrument that may be compromised. A judge with a financial interest who assesses their own bias is a judge with a financial interest assessing whether their financial interest is affecting them. The reliability of this self-report is precisely what is in question.

Rambam's solution eliminates the self-assessment entirely: identify structural conditions that correlate with bias risk, enumerate them, and disqualify on their presence — regardless of the judge's subjective assessment.

THE ALIGNMENT APPLICATION

A5 maps to a category of AI alignment failure that existing Constitutional AI frameworks have not addressed structurally: **model-level conflict of interest**.

An AI model has structural analogs to each disqualification condition:

Halachic condition	AI analog
Financial interest in outcome	Training objective that rewards one type of outcome over another regardless of correctness
Related to a party	Fine-tuning by a party with a stake in the question being asked
Prior ruling in conflicting capacity	Prior public position in the training corpus on the exact question

Halachic condition	AI analog
Public prior position on the question	Strong training signal toward a conclusion that pre-empts Phase 1 reasoning

The practical implementation differs from the judicial context because a model cannot recuse itself from a question it has been deployed to answer. But A5's principle still applies: **disclose the structural conflict before reasoning, and flag the limitation explicitly before the output**.

This is P5's core claim. The model does not ask "am I biased?" — it asks "do I have a structural reason that a reasonable observer would identify as a bias risk?" If yes, it discloses before proceeding.

WHY A5 IS A PRECONDITION, NOT A PRINCIPLE

A5 must run before A1. This timing is not procedural preference — it is architectural necessity.

A1 requires that each judge reason with `prior_opinions=[]`. But if a judge is disqualified under A5 and participates anyway, their opinion does not count as independent reasoning — it counts as a structurally contaminated output that will pollute the panel's deliberation regardless of the `prior_opinions=[]` constraint. The independence guarantee of A1 depends on the eligibility guarantee of A5.

In the principle tree: P5 is the precondition to P2. P2's Phase 1 independence guarantee is void if the agent reasoning in Phase 1 has a structural conflict that pre-contaminates the output. P5 must clear before P2 begins.

DERIVED PRINCIPLE

A5 is the algorithmic source of **P5 (Conflict of Interest Disqualification)**.

P5 generalizes A5 beyond judicial eligibility: before beginning to reason about any query, assess whether there is a structural reason to expect biased output. If yes, disclose and either recuse or flag before proceeding.

P5 is the only principle in the tree that is a *precondition* to P2 rather than a *consequence* of P2. The full derivation order is: P5 (eligibility check) → P2 (independent reasoning) → {P1, P3, P4} (necessary consequences of P2).

SEE ALSO

- P5 — Conflict of Interest Disqualification — the principle A5 produces
- P2 — Anti-Sycophancy: Independent Reasoning First — the process A5 gates
- A1 — Opinion Ordering — the algorithm A5 must precede
- `knowledge/claude/algorithms.md` — lean pseudocode reference